

Q2. First of all we updated the `proc.h` file where we added **ctime**, **retime**, **rutime** and **stime** inside the `proc` struct. After that we defined the functions **statistics_handler**(Will update the process time values depending upon it's state) and **wait2**(Extends the wait system call) inside the same file.

```

38 struct proc {
39     uint sz;           // Size of process memory (bytes)
40     pde_t* pgdir;      // Page table
41     char *kstack;      // Bottom of kernel stack for this process
42     enum procstate state; // Process state
43     int pid;           // Process ID
44     struct proc *parent; // Parent process
45     struct trapframe *tf; // Trap frame for current syscall
46     struct context *context; // switch() here to run process
47     void *chan;         // If non-zero, sleeping on chan
48     int killed;         // If non-zero, have been killed
49     struct file *ofile[NOFILE]; // Open files
50     struct inode *cwd;  // Current directory
51     char name[16];      // Process name (debugging)
52     int ctime;          // Time of creation
53     int stime;          // Number of clock ticks during which the process was waiting for I/O
54     int retime;         // Number of clock ticks during which the process waited
55     int rutime;         // Number of clock ticks during which the process was running
56 };
57 // Process memory is laid out contiguously, low addresses first:
58 // text
59 // original data and bss
60 // fixed-size stack
61 // expandable heap
62
63 void statistics_handler(); // Will update the process time values depending upon it's state
64 int wait2(int*, int*, int*); // Extends the wait system call

```

proc.h

Then we updated the `proc.c` file where we implemented the **wait2** and **statistics_handler** functions. Also we updated the **allocproc** function.

int wait2(int* retime, int* rutime, int* stime):-

This function takes 3 arguments i.e. `retime`, `rutime` and `stime` and updates these values inside it. It returns the PID of the successfully terminated child process if it is not possible then it will return 1. In this function we are first checking where the current process is the parent of the process currently being searched. If it is so then we are checking whether the process is in ZOMBIE or not. If it is then we are terminating that process and assigning corresponding time statistics values inside the arguments taken and then finally return the PID of the process.

```

329 int wait2(int *retime, int *rutime, int *stime)
330 {
331     struct proc *p;
332     int havekids, pid;
333     struct proc *curproc = myproc();
334
335     acquire(&table.lock);
336     for (;;)
337     {
338         // Scan through table looking for exited children.
339         havekids = 0;
340         for (p = table.proc; p < &table.proc[NPROC]; p++)
341         {
342             if (p->parent != curproc)
343                 continue;
344             havekids = 1;
345             if (p->state == ZOMBIE)
346             {
347                 // Found one.
348                 pid = p->pid;
349                 *retime = p->retime;
350                 *stime = p->stime;
351                 *rutime = p->rutime;
352
353                 kfree(p->kstack);
354                 p->kstack = 0;
355                 freevm(p->pgdir);
356                 p->pgdir = 0;
357                 p->parent = 0;
358                 p->name[0] = 0;
359                 p->killed = 0;
360
361                 p->ctime = 0;
362                 p->retime = 0;
363                 p->stime = 0;
364                 p->rutime = 0;
365
366                 p->state = UNUSED;
367                 release(&table.lock);
368                 return pid;
369             }
370         }
371
372         // No point waiting if we don't have any children.
373         if (!havekids || curproc->killed)
374         {
375             release(&table.lock);
376             return 1;
377         }
378
379         // Wait for children to exit. (See wakeup call in proc_exit.)
380         sleep(curproc, &table.lock); //DOC: wait-sleep
381     }
382 }

```

void statistics_handler() :-

In this function we will iterate through each process present inside the process table and will update the process according to the state(RUNNING, SLEEPING, RUNNABLE) in which it is present.

```

void statistics_handler()
{
    struct proc *p;
    acquire(&ptable.lock);
    for (p = ptable.proc; p < &ptable.proc[NPROC]; p++)
    {
        switch (p->state)
        {
            case SLEEPING:
                p->stime++;
                break;
            case RUNNABLE:
                p->retime++;
                break;
            case RUNNING:
                p->rtime++;
                break;
            default:;
        }
    }
    release(&ptable.lock);
}

```

static struct proc* allocproc(void)

Here as a new process is getting created therefore we are assigning **ctime=ticks** and **retime**, **rtime**, **stime** as zero.

```

static struct proc *
allocproc(void)
{
    struct proc *p;
    char *sp;

    acquire(&ptable.lock);

    for (p = ptable.proc; p < &ptable.proc[NPROC]; p++)
        if (p->state == UNUSED)
            goto found;

    release(&ptable.lock);
    return 0;

found:
    p->state = EMBRYO;
    p->pid = nextpid++;

    p->ctime = ticks;
    p->rtime = 0;
    p->retime = 0;
    p->stime = 0;
    release(&ptable.lock);

    // Allocate kernel stack.
    if ((p->kstack = kalloc()) == 0)
    {
        p->state = UNUSED;
        return 0;
    }
    sp = p->kstack + KSTACKSIZE;
}

```

Then we updated the **trap.c** file. Inside the trap function we called the **statistics_handler** inside the switch in the case where **tf->trapno** is equal to **T_IRQ + IRQ_TIMER**. Whenever ticks gets updated i.e. incremented we are also calling the **statistics_handler** that time which will update all the processes available.

```

void
trap(struct trapframe *tf)
{
    if(tf->trapno == T_SYSCALL){
        if(myproc()->killed)
            exit();
        myproc()->tf = tf;
        syscall();
        if(myproc()->killed)
            exit();
        return;
    }

    switch(tf->trapno){
    case T_IRQ0 + IRQ_TIMER:
        if(cpuid() == 0){
            acquire(&tickslock);
            ticks++;
            statistics_handler();
            wakeup(&ticks);
            release(&tickslock);
        }
        lapiceoi();
        break;
    case T_IRQ0 + IRQ_IDE:
        ideintr();
        lapiceoi();
        break;
    case T_IRQ0 + IRQ_IDE+1:
        // Bochs generates spurious IDE1 interrupts.
        break;
    case T_IRQ0 + IRQ_KBD:

```

Then I created a system call which will be used to call the wait2 function inside the **sysproc.c** file. In this we are first defining 3 integer pointer variables and placing values inside them (the one received from the user). Then we are calling the wait2 function and then returning the value wait2 returned.

```

93  int sys_wait2(void) {
94      int *retime, *rtime, *stime;
95      if (argptr(0, (void*)&retime, sizeof(retime)) < 0)
96          return -1;
97      if (argptr(1, (void*)&rtime, sizeof(rtime)) < 0)
98          return -1;
99      if (argptr(2, (void*)&stime, sizeof(stime)) < 0)
100          return -1;
101      return wait2(retime, rtime, stime);
102  }
103

```

After that we defined a macro for this function and gave it the system call number 22 inside the **syscall.h** file and this helped us to connect it to the implemented function.

```

C syscall.h > SYS_getpid
1 // System call numbers
2 #define SYS_fork 1
3 #define SYS_exit 2
4 #define SYS_wait 3
5 #define SYS_pipe 4
6 #define SYS_read 5
7 #define SYS_kill 6
8 #define SYS_exec 7
9 #define SYS_fstat 8
10 #define SYS_chdir 9
11 #define SYS_dup 10
12 #define SYS_getpid 11
13 #define SYS_sbrk 12
14 #define SYS_sleep 13
15 #define SYS_uptime 14
16 #define SYS_open 15
17 #define SYS_write 16
18 #define SYS_mknod 17
19 #define SYS_unlink 18
20 #define SYS_link 19
21 #define SYS_mkdir 20
22 #define SYS_close 21
23 #define SYS_wait2 22

```

Then to make this function visible to the entire program I added the line **extern int sys_wait2(void)** inside the **syscall.c** file and thus it connects the shell and the kernel. Then inside the syscalls array we added sys_wait2 at 22nd position.

```

101 extern int sys_sleep(void);
102 extern int sys_unlink(void);
103 extern int sys_wait(void);
104 extern int sys_write(void);
105 extern int sys_uptime(void);
106 extern int sys_wait2(void);
107
108 static int (*syscalls[])(void) = {
109 [SYS_fork] sys_fork,
110 [SYS_exit] sys_exit,
111 [SYS_wait] sys_wait,
112 [SYS_pipe] sys_pipe,
113 [SYS_read] sys_read,
114 [SYS_kill] sys_kill,
115 [SYS_exec] sys_exec,
116 [SYS_fstat] sys_fstat,
117 [SYS_chdir] sys_chdir,
118 [SYS_dup] sys_dup,
119 [SYS_getpid] sys_getpid,
120 [SYS_sbrk] sys_sbrk,
121 [SYS_sleep] sys_sleep,
122 [SYS_uptime] sys_uptime,
123 [SYS_open] sys_open,
124 [SYS_write] sys_write,
125 [SYS_mknod] sys_mknod,
126 [SYS_unlink] sys_unlink,
127 [SYS_link] sys_link,
128 [SYS_mkdir] sys_mkdir,
129 [SYS_close] sys_close,
130 [SYS_wait2] sys_wait2,
131 };

```

Then to connect the user's call to the system call function we created a user level system call definition inside the **usys.S** file.

```
usys.S
1  #include "syscall.h"
2  #include "traps.h"
3
4  #define SYSCALL(name) \
5      .globl name; \
6      name: \
7          movl $SYS_ ## name, %eax; \
8          int $T_SYSCALL; \
9          ret
10
11 SYSCALL(fork)
12 SYSCALL(exit)
13 SYSCALL(wait)
14 SYSCALL(pipe)
15 SYSCALL(read)
16 SYSCALL(write)
17 SYSCALL(close)
18 SYSCALL(kill)
19 SYSCALL(exec)
20 SYSCALL(open)
21 SYSCALL(mknod)
22 SYSCALL(unlink)
23 SYSCALL(fstat)
24 SYSCALL(link)
25 SYSCALL(mkdir)
26 SYSCALL(chdir)
27 SYSCALL(dup)
28 SYSCALL(getpid)
29 SYSCALL(sbrk)
30 SYSCALL(sleep)
31 SYSCALL(uptime)
32 SYSCALL(wait2)]
```

Then comes the final part, inside the **user.h** (user header file) I added the user call **int wait2(int*, int*, int*)** . And this user call is the one which takes 3 arguments retime, rutime, stime and updates them and then finally returns the PID of the terminated child process.

```
// system calls
int fork(void);
int exit(void) __attribute__((noreturn));
int wait(void);
int pipe(int*);
int write(int, const void*, int);
int read(int, void*, int);
int close(int);
int kill(int);
int exec(char*, char**);
int open(const char*, int);
int mknod(const char*, short, short);
int unlink(const char*);
int fstat(int fd, struct stat*);
int link(const char*, const char*);
int mkdir(const char*);
int chdir(const char*);
int dup(int);
int getpid(void);
char* sbrk(int);
int sleep(int);
int uptime(void);
int wait2(int *,int *,int *);
```

Then we created a user level test file named **wait2_test.c** which calls this function and prints the respective information. In this file we are first creating new duplicate child processes by using **fork()** and then calling **wait2** and then storing it's value in the **pid** variable. Then if the **fork** returns 0 value then it means that the current process is a child process. Therefore for that iteration we are just

printing "Child Process" in the console. If it is not zero then we will check whether pid==fork's return value; if it so then we will print the retime,rutime and stime as it returns the terminated child process. Else we will print the error that is not able to terminate the child process on the console.

```
C wait2_test.c > main()
1  #include "types.h"
2  #include "user.h"
3
4  int main()
5  {
6      int retime, rutime, stime, n1;
7      for (int a = 0; a < 3; a++)
8      {
9          n1 = fork();
10         int pid = wait2(&retime, &rutime, &stime);
11         if (!n1)
12         {
13             printf(1, "Child Process.\n");
14             continue;
15         }
16         if (pid == n1)
17         {
18             printf(1, "| PID: %d | Retime: %d | Rutime: %d | Stime: %d |\n", n1, retime, rutime, stime);
19         }
20         else
21         {
22             printf(1, "Not able to terminate the child process.\n");
23         }
24     }
25     exit();
26 }
27
```

```
$ wait2_test
Child Process.
Child Process.
Child Process.
| PID: 6 | Retime: 0 | Rutime: 0 | Stime: 0 |
| PID: 5 | Retime: 0 | Rutime: 3 | Stime: 0 |
Child Process.
| PID: 7 | Retime: 0 | Rutime: 1 | Stime: 0 |
| PID: 4 | Retime: 0 | Rutime: 4 | Stime: 4 |
Child Process.
Child Process.
| PID: 9 | Retime: 0 | Rutime: 1 | Stime: 0 |
| PID: 8 | Retime: 0 | Rutime: 4 | Stime: 1 |
Child Process.
| PID: 10 | Retime: 0 | Rutime: 1 | Stime: 0 |
$
```

The above picture depicts the output of the file. As we can see there are 7 child processes present and therefore we can terminate 7 child processes whose statistics are also shown.

So to make this program available for xv6 source code for compilation we included wait2_test in the User programs or UPROGS by adding the line `_wait2_test\` inside the **UPROGS** present inside the **Makefile**.

Then we reloaded the QEMU terminal by executing following commands:-

- 1.) make clean
- 2.) make qemu

These commands will clean, compile and link things respectively.

Then I executed the **wait2_test** to get the desired output. We can also list all the **fs.img** content using ls command through which we ensure that the wait2_test command is available.